

Model Optimization and Tuning Phase Report

Date	20 June 2025
Team ID	SWTID1749791625
Project Title	Smart Lender- Applicant Credibility Prediction for Loan Approval
Maximum Marks	10 Marks

Model Optimization and Tuning Phase


The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.




Hyperparameter Tuning Documentation (6 Marks):

Model	Tuned Hyperparameters	Optimal Values
Decision Tree	<pre># Define the Decision Tree classifier dt_classifier = DecisionTreeClassifier() # Define the hyperparameters and their possible values for tuning param_grid = { 'criterion': ['gini', 'entropy'], 'splitter': ['best', 'random'], 'max_depth': [None, 10, 20, 30, 40, 50], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4], }</pre>	<pre># Evaluate the performance of the tuned model accuracy = accuracy_score(y_test, y_pred) print("Optimal Hyperparameters:", best_params) print("Accuracy on Test Set:", accuracy) Optimal Hyperparameters: {'metric': 'euclidean', 'n_neighbors': 9, 'weights': 'uniform'} Accuracy on Test Set: 0.8536585365853658</pre>
Random Forest	<pre>from sklearn.ensemble import RandomForestClassifier from sklearn.model_selection import GridSearchCV param_grid = { 'n_estimators': [100, 200, 300], 'max_depth': [None, 10, 20, 30], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4], 'max_features': ['auto', 'sqrt'] } rf = RandomForestClassifier(random_state=42) grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, scoring='accuracy', cv=5, n_jobs=-1, verbose=2) grid_search.fit(X_train_final, y_train_final) print("Best parameters found: ", grid_search.best_params_) print("Best cross-validation accuracy: {:.4f}".format(grid_search.best_score_))</pre>	<pre>0.7605057 0.7621018 0.7621018 0.7621018 0.7621018 0.7621018 0.7623722 0.7623722 0.7623722 0.7623722 0.7623722 0.7623722 Best parameters found: {'max_depth': 20, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 300} Best cross-validation accuracy: 0.7889</pre>

KNN	<pre># Define the KNN classifier knn = KNeighborsClassifier() # Define the hyperparameters and their possible values for tuning param_grid = { 'n_neighbors': [3, 5, 7, 9], 'weights': ['uniform', 'distance'], 'metric': ['euclidean', 'manhattan'] } # Define the KNN classifier knn = KNeighborsClassifier() # Define the hyperparameters and their possible values for tuning param_grid = { 'n_neighbors': [3, 5, 7, 9], 'weights': ['uniform', 'distance'], 'metric': ['euclidean', 'manhattan'] }</pre>	<pre># Evaluate the performance of the tuned model accuracy = accuracy_score(y_test, y_pred) print("Optimal Hyperparameters:", best_params) print("Accuracy on Test Set:", accuracy) Optimal Hyperparameters: {'metric': 'euclidean', 'n_neighbors': 9, 'weights': 'uniform'} Accuracy on Test Set: 0.8536585365853658</pre>
Gradient Boosting	<pre>from sklearn.model_selection import GridSearchCV from xgboost import XGBClassifier param_grid = { 'n_estimators': [50, 100, 200], 'max_depth': [3, 5, 7], 'learning_rate': [0.01, 0.1, 0.2], 'subsample': [0.8, 1.0], 'colsample_bytree': [0.8, 1.0] } xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42) grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, scoring='accuracy', cv=5, n_jobs=-1, verbose=2) grid_search.fit(X_train_scaled, y_train)</pre>	<pre>print("Best loading...s found: ", grid_search.best_params_) print("Best cross-validation accuracy: {:.4f}".format(grid_search.best_score_)) Fitting 5 folds for each of 100 candidates, totalling 500 fits /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [18:05:30] WARNING: /workspace/src/learn Parameters: {'use_label_encoder': 0} are not used warnings.warn(msg, UserWarning) Best parameters found: {'colsample_bytree': 1.0, 'learning_rate': 0.01, 'max_depth': 5, 'n_estimators': 200, 's Best cross-validation accuracy: 0.7864</pre>

Performance Metrics Comparison Report (2 Marks):

Model	Optimized Metric																														
Decision Tree	<div><div> Test Set Metrics for Decision Tree:</div><table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>0</td><td>0.46</td><td>0.68</td><td>0.55</td><td>38</td></tr><tr><td>1</td><td>0.82</td><td>0.65</td><td>0.72</td><td>85</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.66</td><td>123</td></tr><tr><td>macro avg</td><td>0.64</td><td>0.67</td><td>0.64</td><td>123</td></tr><tr><td>weighted avg</td><td>0.71</td><td>0.66</td><td>0.67</td><td>123</td></tr></tbody></table></div>		precision	recall	f1-score	support	0	0.46	0.68	0.55	38	1	0.82	0.65	0.72	85	accuracy			0.66	123	macro avg	0.64	0.67	0.64	123	weighted avg	0.71	0.66	0.67	123
		precision	recall	f1-score	support																										
	0	0.46	0.68	0.55	38																										
	1	0.82	0.65	0.72	85																										
	accuracy			0.66	123																										
	macro avg	0.64	0.67	0.64	123																										
	weighted avg	0.71	0.66	0.67	123																										

Random Forest	<div><div> Test Set Metrics for Random Forest:</div><table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>0</td><td>0.52</td><td>0.66</td><td>0.58</td><td>38</td></tr><tr><td>1</td><td>0.83</td><td>0.73</td><td>0.78</td><td>85</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.71</td><td>123</td></tr><tr><td>macro avg</td><td>0.67</td><td>0.69</td><td>0.68</td><td>123</td></tr><tr><td>weighted avg</td><td>0.73</td><td>0.71</td><td>0.72</td><td>123</td></tr></tbody></table></div>		precision	recall	f1-score	support	0	0.52	0.66	0.58	38	1	0.83	0.73	0.78	85	accuracy			0.71	123	macro avg	0.67	0.69	0.68	123	weighted avg	0.73	0.71	0.72	123
	precision	recall	f1-score	support																											
0	0.52	0.66	0.58	38																											
1	0.83	0.73	0.78	85																											
accuracy			0.71	123																											
macro avg	0.67	0.69	0.68	123																											
weighted avg	0.73	0.71	0.72	123																											
KNN	<div><div> Test Set Metrics for KNN:</div><table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>0</td><td>0.54</td><td>0.66</td><td>0.60</td><td>38</td></tr><tr><td>1</td><td>0.83</td><td>0.75</td><td>0.79</td><td>85</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.72</td><td>123</td></tr><tr><td>macro avg</td><td>0.69</td><td>0.71</td><td>0.69</td><td>123</td></tr><tr><td>weighted avg</td><td>0.74</td><td>0.72</td><td>0.73</td><td>123</td></tr></tbody></table></div>		precision	recall	f1-score	support	0	0.54	0.66	0.60	38	1	0.83	0.75	0.79	85	accuracy			0.72	123	macro avg	0.69	0.71	0.69	123	weighted avg	0.74	0.72	0.73	123
	precision	recall	f1-score	support																											
0	0.54	0.66	0.60	38																											
1	0.83	0.75	0.79	85																											
accuracy			0.72	123																											
macro avg	0.69	0.71	0.69	123																											
weighted avg	0.74	0.72	0.73	123																											
Gradient Boosting	<div><div> Test Set Metrics for XGBoost:</div><table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>0</td><td>0.33</td><td>0.79</td><td>0.46</td><td>38</td></tr><tr><td>1</td><td>0.74</td><td>0.27</td><td>0.40</td><td>85</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.43</td><td>123</td></tr><tr><td>macro avg</td><td>0.53</td><td>0.53</td><td>0.43</td><td>123</td></tr><tr><td>weighted avg</td><td>0.61</td><td>0.43</td><td>0.42</td><td>123</td></tr></tbody></table></div>		precision	recall	f1-score	support	0	0.33	0.79	0.46	38	1	0.74	0.27	0.40	85	accuracy			0.43	123	macro avg	0.53	0.53	0.43	123	weighted avg	0.61	0.43	0.42	123
	precision	recall	f1-score	support																											
0	0.33	0.79	0.46	38																											
1	0.74	0.27	0.40	85																											
accuracy			0.43	123																											
macro avg	0.53	0.53	0.43	123																											
weighted avg	0.61	0.43	0.42	123																											

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Ensemble Voting Classifier	After training and evaluating multiple machine learning models — including Decision Tree, Random Forest, K-Nearest Neighbors (KNN), and XGBoost — I selected the Ensemble Voting Classifier as the final model. This decision was based on its superior performance across metrics such as accuracy, precision, recall, and f1-score.